

AHMAD FARAZ KHAN

ahmadfk@vt.edu ◇ afkd98.github.io ◇ linkedin.com/in/ahmadfarazkhandurrani/ ◇ +1 540-449-6457

EDUCATION

Virginia Tech - Ph.D. in Computer Science, Blacksburg, VA
Research Focus: Machine Learning Systems

January 2021 – *present*

Virginia Tech - M.S. in Computer Science, Blacksburg, VA

June 2024

LUMS - B.S. in Computer Science, Lahore, Pakistan

May 2020

WORK EXPERIENCE

Google, Mountain View - Software Engineering Intern PhD

Aug 2025 – *present*

Manager: Dr. Shan Li

Foundation models:

- Created a pipeline to generate up to 1 million training data images for image-to-image regression tasks.
- Building generic foundation models for video pre-processing regression tasks, including denoising, super-resolution, smart downscaling, and compression.

Impact: *Contributing to the training of **foundation models** for Youtube video regression.*

IBM Research, Almaden - Research Intern

May 2025 – Aug 2025

Managers: Dr. Taiga Nakamura & Dr. Swanand Ravindra Kadhe

Continual learning & Targeted data generation:

- Designed a post-training self-optimizing loop and domain-specific fine-grained synthetic data generation techniques guided by embedding-based similarity metrics, increasing the training data to accuracy efficiency by $2.5\times$.
- Enabled controlled distributional coverage, continual learning, and autonomous maintenance of foundational models, significantly improving robustness and performance in domain-specific tasks.

Impact: *Submitted **2 patents** and contributed to the training of **IBM's Granite 4.0 model**.*

Virginia Tech - Graduate Research Assistant, DSSL

Spring 2021 – *present*

Advisors: Dr. Ali Butt, Virginia Tech & Dr. Ali Anwar, University of Minnesota

ML Infrastructure & Algorithms Optimization:

- Designed an RLHF approach to fine-tune deep learning compression optimizations without sacrificing accuracy. Increased **resource utilization up to $81\times$** , **scalability by $78\times$** , and **accuracy up to 53%** .
- Developed clustering-based personalized learning solutions for distributed ML systems. Improved the **personalized accuracy by up to 45%** .
- Devised a Direct Preference Optimization (DPO) approach for prompt optimization without separate reward modeling for LLMs; **enhanced score by 27%** over supervised fine-tuning.
- Created a DPO pipeline to mitigate sycophancy, cutting it by **64%** in persona tests and **44%** in preference-driven tests.
- Engineered a context-aware agentic AI DevOps platform for adaptive cloud deployments, reducing human effort and cost by **90%** .
- Created an adaptive aggregator server for collaborative learning with **one million+** nodes; improved **scalability by $4\times$** , **latency by $8\times$** , and **reduced cost by $2\times$** .
- Developed a scheduler balancing efficiency vs. accuracy; improved **accuracy by 57%** and **reduced training time by 40%** .
- Engineered a locality-aware cache for non-training workloads, decreasing **latency and cost by 71%** and **98%** , respectively.

Impact: *Publications at **MLSys'25**, **IPDPS'25**, **ACM EuroSys'24**, **IEEE BigData'24 (Best paper)**, **IEEE CLOUD'22**, **IEEE BigData'22 & 23**, **FL-AAAI'22**, submissions: **TOSEM'26**, **ACL'25**.*

SELECT PUBLICATIONS

“FLStore: Efficient Federated Learning Storage for non-training workloads”, **Ahmad Faraz Khan** et al., *MLSys 2025*.

“IP-FL: Incentive-driven Personalization in Federated Learning”, **Ahmad Faraz Khan** et al., *IPDPS 2025*.

“FLOAT: Federated Learning Optimizations with Automated Tuning”, **Ahmad Faraz Khan** et al., *EuroSys 2024*.

“DynamicFL: Federated Learning with Dynamic Communication Resource Allocation”, Qi Le, Enmao Diao, Xinran Wang, **Ahmad Faraz Khan** et al., *BigData 2024 (Best Paper)*.

“Mitigating Sycophancy in LLMs via Direct Preference Optimization”, Azal Ahmad Khan, Sayan Alam, Xinran Wang, **Ahmad Faraz Khan** et al., *BigData 2024*.

TECHNICAL PROFICIENCY

- Programming Languages:** Python, JavaScript, C++
- Tools and Libraries:** PyTorch, TensorFlow, Hugging Face, LangChain, Ollama, Pandas, SciPy, FLOWER, IBM Federated Learning, Spark MLlib, PySpark, Dask, Hadoop, DeepSpeed, MinIO, AWS Suite, Docker, OpenFaaS, SQL, Kubernetes

ADDITIONAL EXPERIENCES & SERVICES

- Reviewer for COLM 2025; USENIX ATC 2024; Springer Neural Processing Letters 2022–23; IEEE TNSM 2024; PeerJ CS 2024.
- Graduate Teaching Roles:** Web/Cloud Development (Summer 2024 & Fall 2023); Advanced Operating Systems (Spring & Fall 2024); Python Programming (Spring 2020 & Fall 2021); Computer Security (Spring 2022).